

*BASICS*

*Please don't share this to your competitor .*

*I may be wrong .The mother loves all the kids in the world .*

*But when it comes to competition , the mother always wants only her kid to win .*

# *Oracle Architectural Components*

# *Oracle Server*

*Oracle instance + Database*

*Oracle instance*

*MEMORY STRUCTURES*

+

*BACKGROUND PROCESSES*

*User Process*

*Server Process*

*Background process*

# Oracle components

## **Data**

*Actual data*

## **Redo**

*Changes made to the database-Used for recovery*

## **Undo**

*Used for Read Consistency*

# *MEMORY STRUCTURES*



# *MEMORY STRUCTURES*

## *1. PROGRAM GLOBAL AREA*

*Process global area*

## *2. SHARED GLOBAL AREA*

*System global area*

# *Program Global Area*

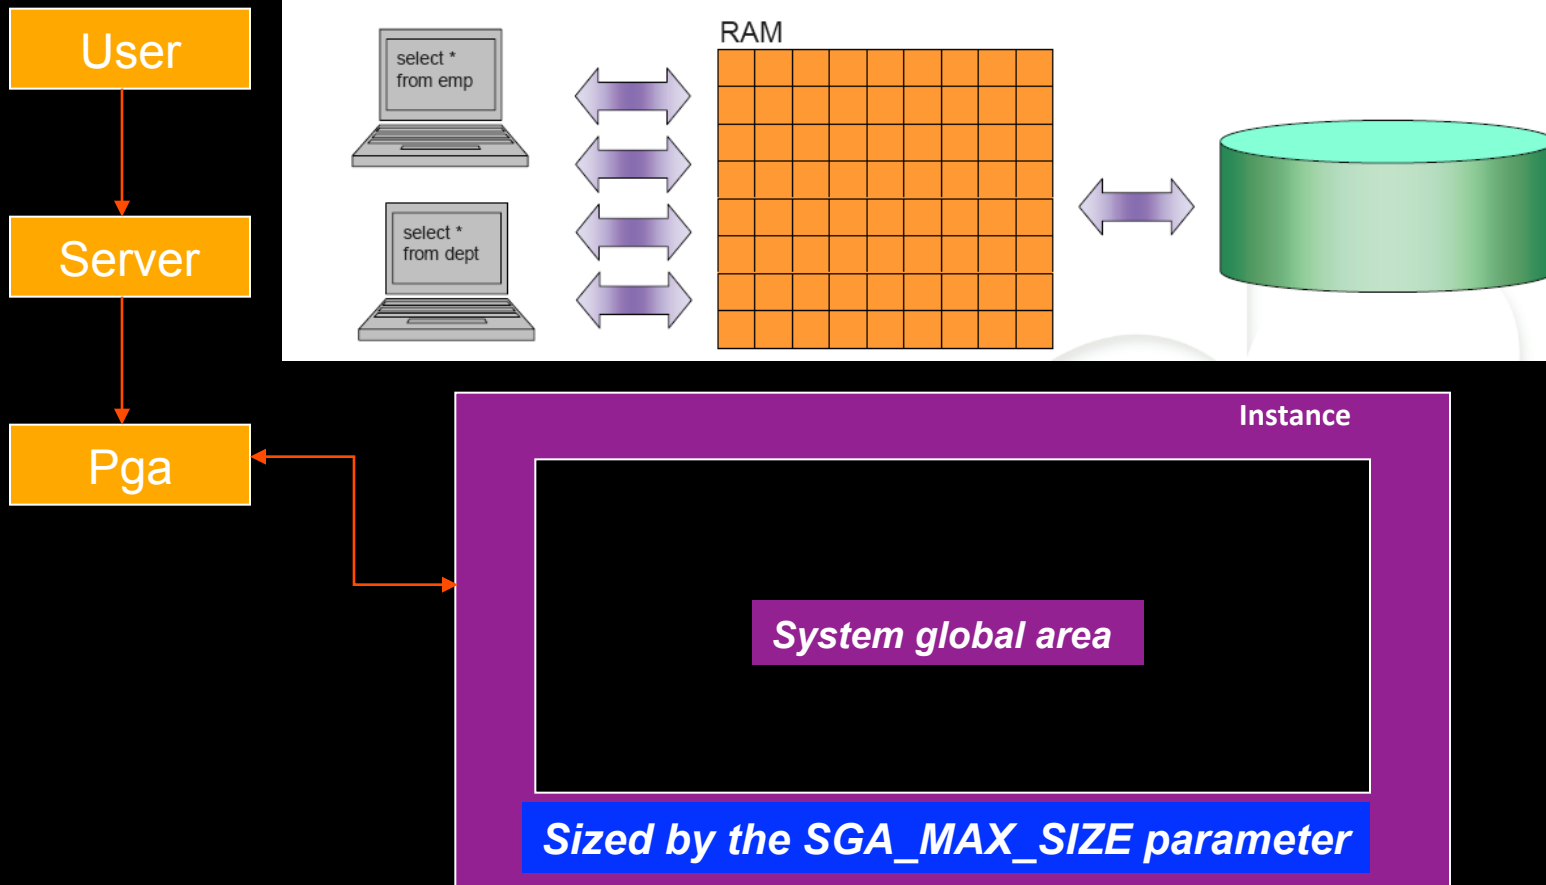
*Memory reserved for each user process connecting to an Oracle database*

*Allocated when a process is created*

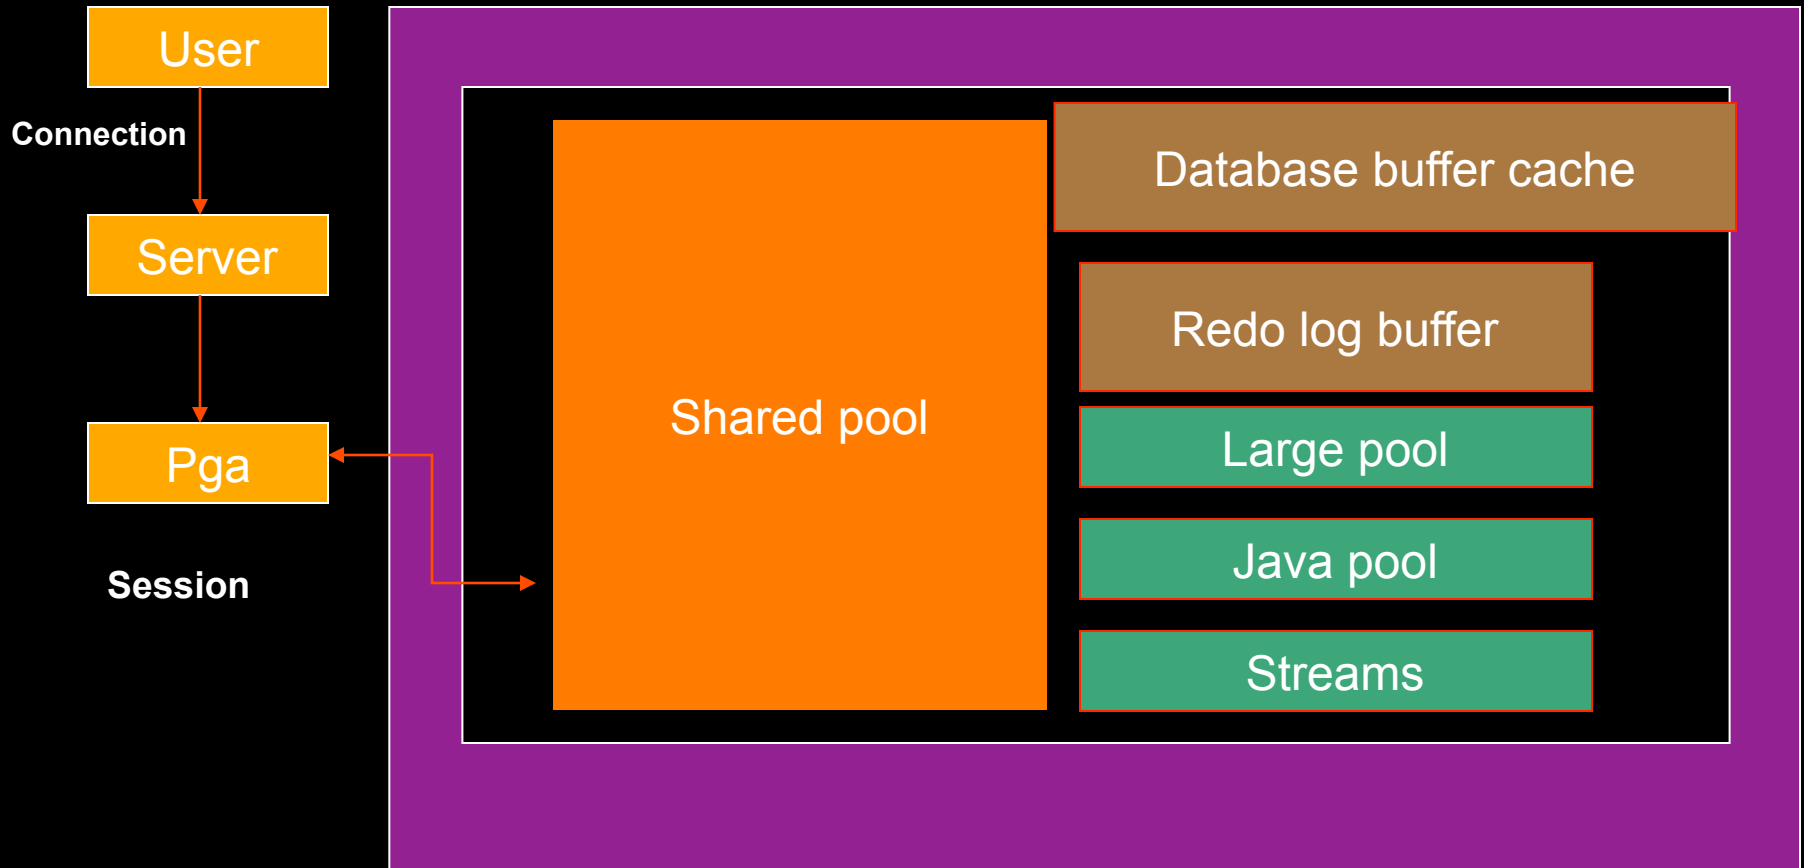
*Deallocated when the process is terminated*

*Used by **only one User process***



# System Global Area



# System Global Area



# Database Buffer Cache




*Before any database block can be used ,  
it must be physically read from disk  
and placed into the database buffer cache*

# Database Buffer Cache

I need File 7, Block 1234


is it already in the cache ?



# *Database Buffer Cache*


File 7, Block 1234 was not in the cache

bring it from disk into the cache



# Database Buffer Cache

Uses *MRU—LRU Mechanism*





# *Database Buffer Cache*

*Database Buffer Cache has 4 different buffers*

*Free Buffer*

*Pinned Buffer*

*DB\_CACHE\_SIZE*



*Dirty Buffer*

*V\$BH*

*Clean Buffer*

*Shared pool*

# Shared pool



*Library cache*

*Data Dictionary cache*

***SHARED\_POOL\_SIZE***

*Helps in Parsing (Only Syntactic)*

*Stores most recently used SQL and PL/SQL statements*

*Enables the sharing of commonly used statements Algorithm*

*Consists of two structures:*

*Shared SQL area*

*Private PL/SQL area*

**V\$LIBRARYCACHE**

# *ROW CACHE*

*Helps in Parsing (Only Semantic)*

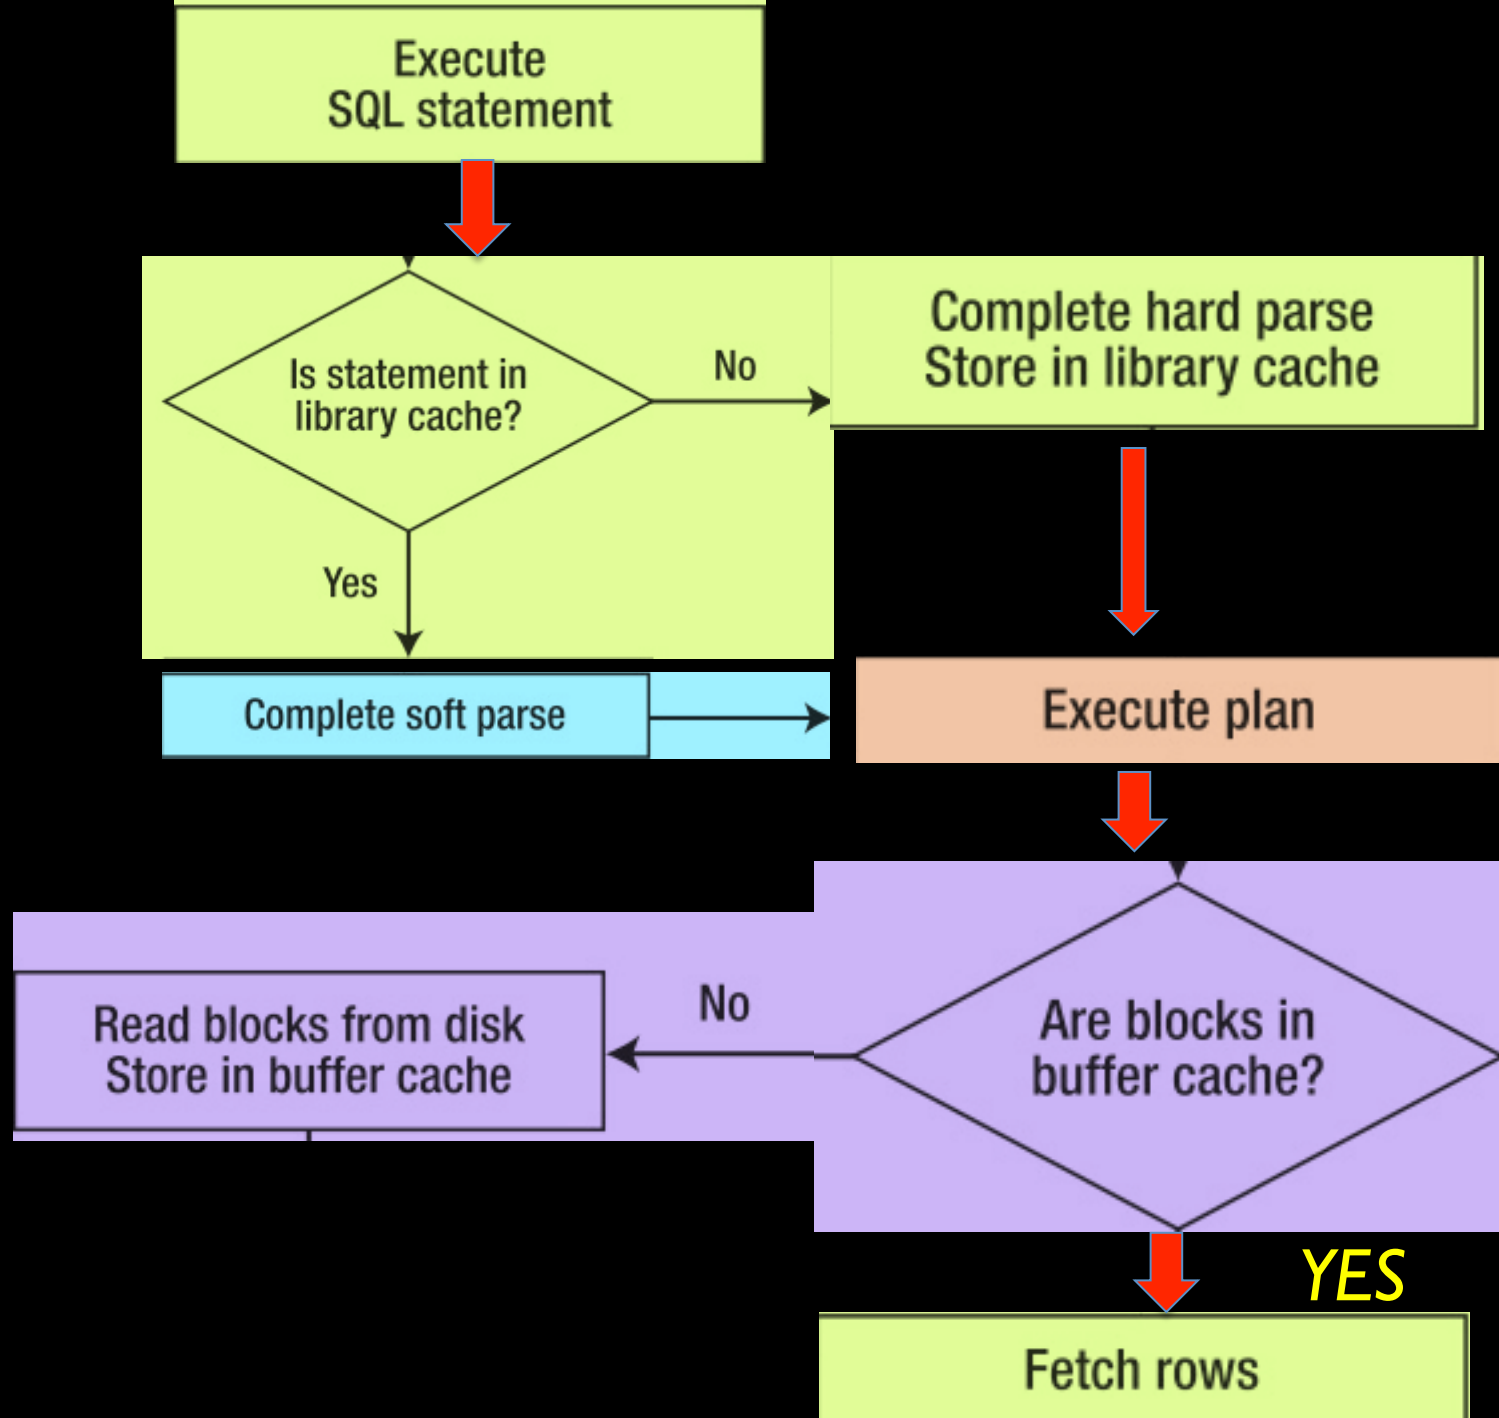
*Collection of the most recently used definitions  
in the database*

*Includes information about database files,  
tables, indexes, columns, users,  
privileges, and other database objects*

**V\$ROWCACHE**

*Steps involved for a sql query*





*Redo log Buffer*

*Records all changes made to the database  
data blocks*

*Primary purpose is recovery*

*Changes recorded within are called redo entries*

*Redo entries contain information to  
reconstruct or redo changes*

# *Background processes*

*Smon*

*Pmon*

*Dbwr*

*Lgwr*

***Ckpt***



*THE REAL HERO*

# *Database*

*How Oracle store Data*

# *Control Files*

*BRAIN OF THE Database*

# *Control File Contents*

*Database name and identifier*

*Names and locations of data files and log files*

*Checkpoint information*

*Data files*



# *Data files*


*Stores actual Data*

*Tablespaces*

*Segments*

*Extents*

*Data blocks*




# Block Utilization Parameters

## Data Block Management

**Two methods are available for managing data blocks:**

- Automatic Segment-Space Management
- Manual Management



INITTRANS  
MAXTRANS  
PCTFREE  
PCTUSED  
FREE LISTS

# *Redo Log Files*

# *How Redo Log Files Work*

*Record all changes made to data*

*Used for recovery*

*Minimum 2 Redo log files are required*

*Redo log files are used in a cyclic fashion.*

# *How Redo Log Files Work*

*When a redo log file is full, LGWR will move to the next log group ,called a log switch*

*Checkpoint operation occurs and  
Information written to the control file*

# *Archive files*

*Archive files are copy of Log files*

*EXPLAIN PLAN / SQL TRACE*



*SET AUTOTRACE ON*

*SET AUTOTRACE TRACEONLY EXPLAIN*

*explain plan for select emp\_id from emp;*

*select \* from table(dbms\_xplan.display);*

## *DBMS\_XPLAN. DISPLAY Examples*

*SELECT \* FROM TABLE*

*(dbms\_xplan.display('PLAN\_TABLE','abc','BASIC'));*

*SELECT \* FROM TABLE*

*(dbms\_xplan.display('PLAN\_TABLE','abc','TYPICAL'))*

*SELECT \* FROM TABLE*

*(dbms\_xplan.display('PLAN\_TABLE','abc','ALL'));*

*BASIC – Displays the minimum information*

*TYPICAL – Displays the relevant information  
in the plan and predicate information ,parallel*

*ALL – All of typical including projections, alias*

*EXPLAIN PLAN*

*set statement\_id = 'X1'*

*FOR SELECT d.deptno , d.dname ,  
COUNT(\*)*

*AS count\_employees ,*

*SUM(e.sal) AS*

*sum\_salaries FROM dept d , emp e*

*WHERE d.deptno = e.deptno*

*GROUP BY d.deptno , d.dname;*

*SELECT plan\_table\_output FROM TABLE*

*( DBMS\_XPLAN.DISPLAY ('PLAN\_TABLE','XI','ALL') )*

*Tkprof*





## Initialisation parameters

- `timed_statistics = TRUE`
- `SQL_TRACE = TRUE`
- `user_dump_dest = '<directory name>'`
- `trace_file_identifer = '<string>'`

## *Enabling and Disabling SQL Trace*

*At the instance level:*

```
SQL_TRACE = TRUE  
ALTER SYSTEM SET SQL_TRACE=FALSE;
```

*At the session level:*

```
Alter session set SQL_TRACE = true  
Alter session set tracefile_identifler= 'sense'
```

# TKPROF performance



TKPROF --Converts the trace file into a readable format.

TKPROF <trace filename> <output filename>

# TKPROF



Set auto trace on

call	count	cpu	elapsed	disk	query	current	rows
-----	-----	-----	-----	-----	-----	-----	-----
Parse	1	0.00	0.00	0	0	0	0
Execute	3117	0.83	0.92	0	0	0	3113
Fetch	3117	7.03	9.88	40	55902	0	3105
-----	-----	-----	-----	-----	-----	-----	-----
total	6235	7.86	10.80	40	55902	0	6218

# TKPROF output

```
*****
*
count      = number of times OCI procedure was executed

cpu        = cpu time in seconds executing

elapsed    = elapsed time in seconds executing

disk       = number of physical reads of buffers from disk

query      = number of buffers gotten for consistent read

current    = number of buffers gotten in current mode (usually for update)

rows       = number of rows processed by the fetch or execute call
*****
*
```

# *Explain plan*

set autotrace on

Set autotrace on explain

Set autotrace on statistics

Set autotrace traceonly explain

Set autotrace traceonly statistics

Set autotrace traceonly

Set autotrace off

# V\$views

## User/Session

V\$LOCK

V\$OPEN\_CURSOR

V\$PROCESS

V\$SORT\_USAGE

V\$SESSIONT/P

V\$SESSTAT T/P

V\$TRANSACTION

V\$SESSION\_EVENT

V\$SESSION\_WAIT

*V\$SQL*

*V\$SQLAREA*

*V\$WAITSTAT*

*V\$SESSION*

*V\$SESS\_IO*

*V\$SQL\_PLAN*

*V\$SQL\_PLAN\_STATISTICS*

*V\$OPEN\_CURSOR*

*V\$SESSION\_LONGOPS*

*V\$SYSSTAT*

*V\$SESSION\_EVENT*

*V\$SESSION\_WAIT*

*V\$SYSTEM\_EVENT*


*V\$FILESTAT*

*Indexes*



# *Access paths*

- *Full table scan*
- *Sample scans*
- *Rowid scans*
- *Index full scans*
- *Index unique scan*
- *Index range scan*
- *Index fast full scans*
- *Index joins*
- *Bitmap joins*
- *Cluster scans*
- *Hash scans*



## *So what's a full-table scan?*

- Process of reading all blocks of a database table sequentially*
- Full-table scans may cause of bad query performance*

# *Why ---Full table scan*

*Missing indexes on large tables*

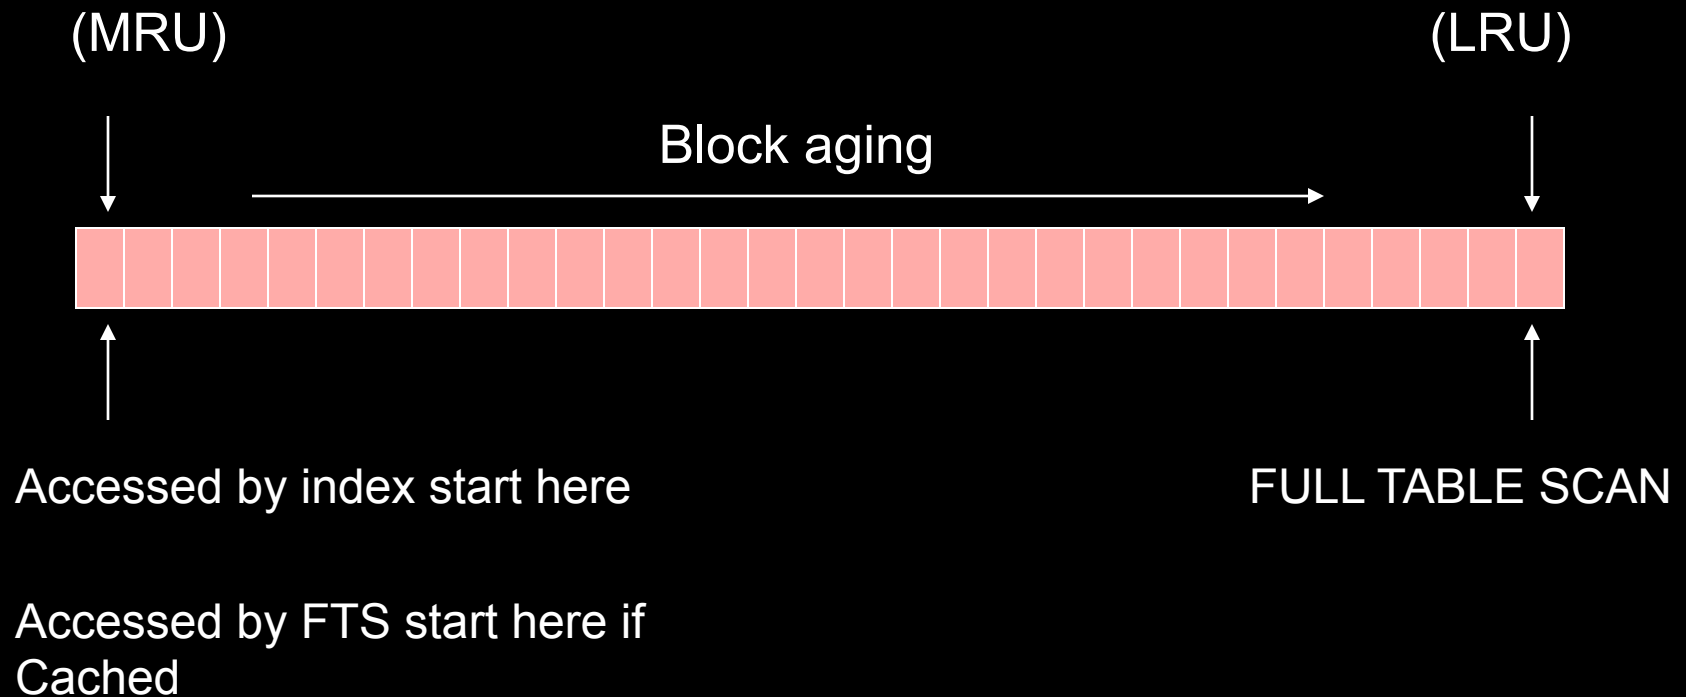
*A poor indexing strategy on large tables*

*Missing table or index statistics*

*Why Full table scan is bad*

# *FTS effects on buffer cache*

## Buffer Cache LRU List



*INDEX*

*ACCESS*

*Last few pages of any book*

1. *B\* tree Index*
2. *Bitmap Index*
3. *Reverse key Index*
4. *Function Based Index*
5. *Bitmap join Index*
6. *Partitioned Index*
7. *Virtual Index*
- 8.



- *Invisible Indexes*



# *Concept of an Index*

*An index contains a pointer for each table row (rowid)*

*To access the data rows quickly*

*Create index emp\_idx (eno) ;*

*Create index emp\_idx (eno,ename,sal) ;*

*Create unique index emp\_idx (eno);*

# *Different Index scans*


<i>Unique scan</i>	<i>1 record</i>
<i>Range Scan</i>	<i>More than 1 record</i>
<i>skip scan</i>	
<i>fast full scan</i>	<i>count (*)</i>
<i>index-join</i>	<i>more than 1 table</i>

*How Oracle store Index data*

# Insertion

Inserting 'ENG','SCO' and 'USA'

Before



After




# Insertion

Inserting 'BEL'

Before



After



# *BITMAP INDEX*

*For column that has a low cardinality  
( few distinct values)*

<i>Examples</i>	<i>Male / female</i>
	<i>Yes / No</i>
	<i>Good / bad</i>
	<i>Colours</i>


# *BITMAP INDEX*

*Create bitmap index emp\_idx (Gender) ;*



# Bitmap Indexes

*Select \* from emp WHERE Gender = 'M' and Color = 'RED';*



*Red: <0,1,0,1,0,1>*

*Mal <1,0,1,1,0,1>*

# REVERSE Key Indexes

KEY	ROWID
-----	-----
1257	0000000F.0002.0001
<b>2877</b>	0000000F.0006.0001
4567	0000000F.0004.0001
<b>6657</b>	0000000F.0003.0001
8967	0000000F.0005.0001
9637	0000000F.0001.0001
9947	0000000F.0000.0001
...	...

EMPLOYEE_ID	LAST_NAME ...
-----	-----
7499	ALLEN
7369	SMITH
7521	WARD ...
<b>7566</b>	JONES
7654	MARTIN
7698	BLAKE
<b>7782</b>	CLARK

*create unique index i1\_t1 ON t1(c1) REVERSE ;*

*alter index i2\_t1 REBUILD REVERSE ;*

## *Function based index:*

*Create index emp\_idx on emp  
(UPPER(ename));*

# Finding index usage

---

I

```
select  
index_name,monitoring,used,start_monitoring,  
end_monitoring from v$object_usage;
```

# *Finding index usage*

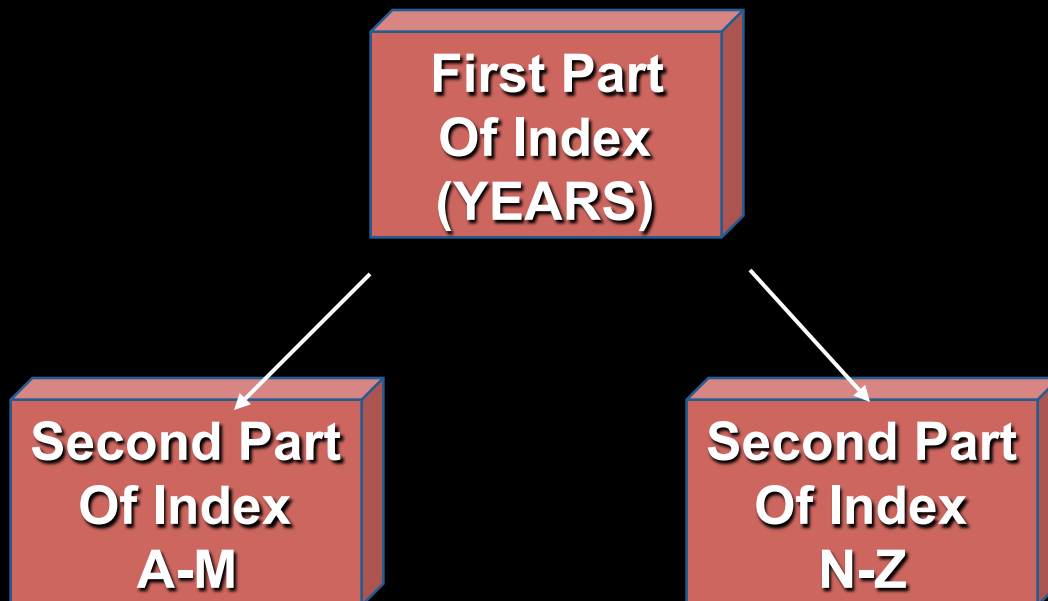
*ALTER INDEX <index> MONITORING USAGE.;*

*ALTER INDEX <index> NOMONITORING usage ;*

*Index Skip Scan*

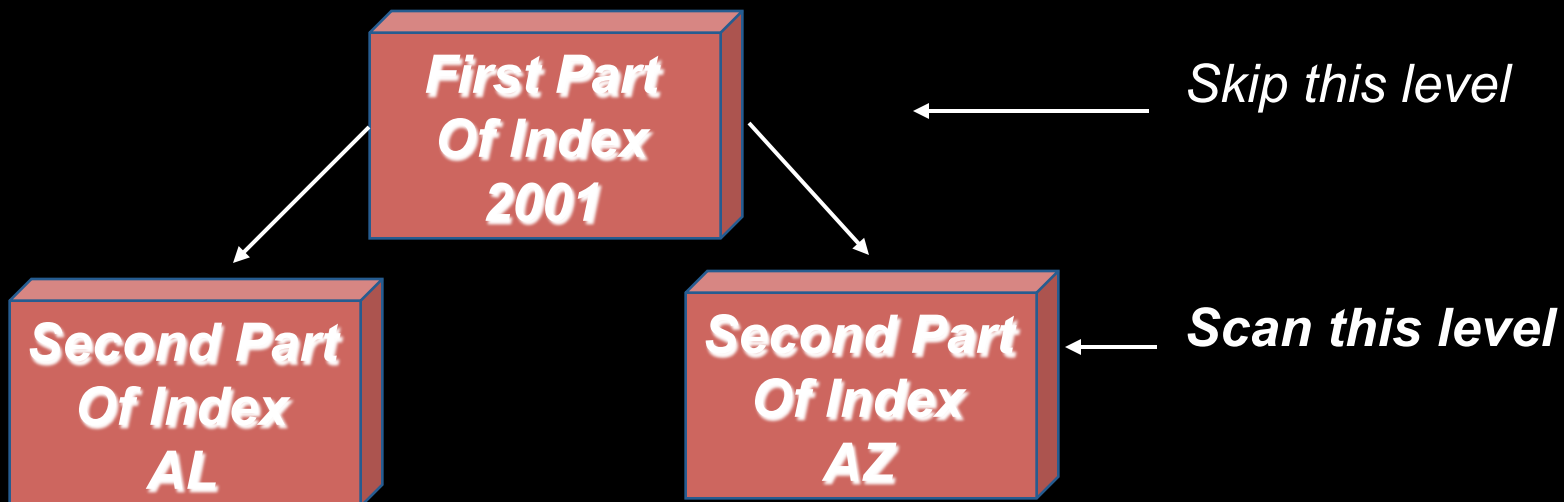
# Index Skip Scan

*Create index year\_state\_idx on test2(year, state);*



# Index Skip Scan

*SELECT COUNT (\*) FROM TEST2  
WHERE STATE= 'AL'*





```
create index IX_SS on EMP(DEPTNO,SAL);
```

```
select /*+ index_ss(EMP IX_SS) */ * from  
emp where SAL < 1500;
```

# *Views for Indexes*

*User\_indexes*

*User\_ind\_columns*

*V\$object\_usage*

*Index\_stats*

# *Index Maintenance*

---

*Alter index emp\_idx rebuild ;*

***WHAT IS***

***REBUILD***

***Removing***

***Un necessary***

---

*Why Should I Rebuild the Index*

# *Why should I rebuild Index*

*Oracle does not delete **Index entries** of **Records** that are deleted by the user.*

*So we rebuild ,to reclaim the space occupied by the deleted rows*

*Performance will become worse as oracle may choose **Full table scan** instead of Index scan*

# *Benefits of Rebuilding Index*

---

*Space occupied by the deleted rows is reclaimed*

*Performance will become better as Oracle  
Chooses Index scan after rebuild .*



# *When Should I Rebuild Index*

---

*Height is  $> 4$*

*Del\_lf\_rows is  $> 20\%$*

*Poor clustering factor*

# *Checking Clustering Factor*

---

*Execute dbms\_stats.gather\_table\_stats( scott,  
'EMP', cascade=>true );*

*Cascade is*

*Gather\_table\_stats + Gather\_index\_stats*

*In a single query*

*Select blocks, num\_rows from  
user\_tables*

*Where table\_name = 'BTREETABLE';*

---

*Select index\_name, blevel, clustering\_factor  
from user\_indexes  
where table\_name = 'BTREETABLE'*

*Blevel ---Btree level*

*“Good” CF — If CF is closer to **Blocks** in table*

*“Bad” CF — If CF is closer to **Rows** in table*

# *Bitmap Join Index*

```
CREATE BITMAP INDEX empdept_bji  
ON employees(d.deptno)  
FROM employees e, departments d  
WHERE d.deptno = e.deptno;
```

# *Invisible Indexes*

```
CREATE INDEX ind1 ON luck(object_id)  
INVISIBLE;
```

```
ALTER INDEX ind1 INVISIBLE ;
```

*Optimizer does not consider this index:*



*Optimizer can consider this index:*

*ALTER INDEX ind1 VISIBLE;*

*Select /\*+ index(luck ind1) \*/ \* from  
luck where object\_id < 1000;*

# Index Organised Tables

## Heap-organised table


```
CREATE TABLE team  
(  
    team_key VARCHAR2(3),  
    team_name  
    VARCHAR2(50),  
    country_key  
    VARCHAR2(3)  
    CONSTRAINT team_pk  
    PRIMARY KEY  
(team_key);  
)  
ORGANIZATION HEAP;
```

## Index-organised table

```
CREATE TABLE team  
(  
    team_key VARCHAR2(3),  
    team_name  
    VARCHAR2(50),  
    country_key  
    VARCHAR2(3)  
    CONSTRAINT team_pk  
    PRIMARY KEY  
(team_key);  
)  
ORGANIZATION INDEX;
```


# *Index Organised Tables*

*Heap-organised table*



*Index-organised table*

*Root*  
*Branch*  
*Leaf*  
*Table*



# *Index Organised Tables*

*You need to have primary key*

*Don't use IOT if data volume is high*

*You can also create indexes for other columns*

*Don't create more than 2 indexes on IOT*

# *Partitioned Indexes*

## *Local Indexes:*

*single index is created of the partitioned table.*

## *Global Indexes:*

*– Single Index is created For all Partitions*

## *Partitioned Indexes examples*

*create index idx\_local on emp(id) local ;*

*create index idx\_global on emp(id) global ;*

## *Virtual Indexes examples*


*create index emp\_idx on emp (object\_id)  
nosegment;*

*alter session set "\_use\_nosegment\_indexes"=true;*

*Does not occupy space in the Database*

# Index Clustering Factor

Must read all blocks to retrieve all As



*POOR CF*

*GOOD CF*

Only need to read one block to retrieve all As



*“Good” CF generally has value closer to blocks in table*

*“Bad” CF generally has a value closer to rows in table –Oracle does not use index*

# *What to do in case of poor CF*

*Create Histograms:*

***Histogram** Stores Data distribution information*

```
exec dbms_stats.gather_table_stats  
(user,'t1',method_opt=>'for all columns size 1' ,  
cascade=>TRUE);
```

*Always specify more than 1 .At least 10*

*Solution for Poor Clustering Factor*

*Increase freelists (not more than 12)*

*Your tablespace should be in MSSM*

*Optimizer*

***Why Do You Need an Optimizer?***

*Find the most efficient mechanism  
for executing any SQL statement*

*JOB OF OPTIMIZER*

*Two main components:*

*Query Transformations*

*Access Path Selection*



*Example:*

*select \* from dept  
where deptno in (select deptno from emp  
where job = 'CLERK')*

*select d.\* from dept d,  
(select distinct deptno from emp where job  
= 'CLERK') e  
where d.deptno = e.deptno*

# *Access Path Selection*

*For each table, choose the access path (table scan, index scan, etc)*

*For each join, choose the join method (nested-loop, sort-merge, hash, etc)*

*Choose the join order for the tables*

*Takes care of*

*Selectivity*

*Cardinality*

*COST*

# Optimizer Parameters

*CURSOR\_SHARING: SIMILAR, EXACT, FORCE*

*DB\_FILE\_MULTIBLOCK\_READ\_COUNT*

*PGA\_AGGREGATE\_TARGET*

*STAR\_TRANSFORMATION\_ENABLED*

*RESULT\_CACHE\_MODE: MANUAL, FORCE*

*RESULT\_CACHE\_MAX\_SIZE*

*RESULT\_CACHE\_MAX\_RESULT*

*RESULT\_CACHE\_REMOTE\_EXPIRATION*

*OPTIMIZER\_INDEX\_CACHING*

*OPTIMIZER\_INDEX\_COST\_ADJ*

# *Optimizer Parameters*

*OPTIMIZER\_FEATURES\_ENABLED*

*OPTIMIZER\_MODE:ALL\_ROWS, FIRST\_ROWS,  
FIRST\_ROWS\_n*

*OPTIMIZER\_CAPTURE\_SQL\_PLAN\_BASELINES*

*OPTIMIZER\_USE\_SQL\_PLAN\_BASELINES*

*OPTIMIZER\_DYNAMIC\_SAMPLING*

*OPTIMIZER\_USE\_INVISIBLE\_INDEXES*

*OPTIMIZER\_USE\_PENDING\_STATISTICS*

# *Gathering Statistics*

*Used to determine:*

*Table access cost*

*Join cardinality*

*Join order*

## *Table Statistics:*

*Number of rows & blocks & empty blocks*

*Average row length*

*Average free space per block*

*(Number of chained rows (CHAIN\_CNT))*



# *Index Statistics:*

*B\*-tree level*

*Distinct keys*

*No of leaf blocks*

*Clustering factor*

*(AVG\_LEAF\_BLOCKS\_PER\_KEY)*

*(AVG\_DATA\_BLOCKS\_PER\_KEY)*

# *Column statistics*

*No of distinct values, no of nulls,  
average length,  
min, max,*

*Histograms*

*(data distribution when the column data is  
skewed)*

*user\_tables*

*user\_all\_tables*

*user\_indexes*

*user\_XXX\_partitions*

*user\_XXX\_subpartitions*

*user\_tab\_cols*

*user\_tab\_col\_statistics*

*user\_(sub)part\_col\_statistics*

*GATHER\_TABLE\_STATS*

*GATHER\_INDEX\_STATS*

*GATHER\_SCHEMA\_STATS*

---

*Execute dbms\_stats.gather\_table\_stats( scott,  
'EMP', cascade=>true );*

*Execute dbms\_stats.gather\_index\_stats  
( winner, 'emp\_idx' );*

*Execute dbms\_stats.gather\_schema\_stats  
(ownname => 'WINNER');*

# Gathering statistics for Partitioned Table

```
exec dbms_stats.gather_table_stats  
( owner=>'winner', Tabname=>'emp',  
partname=>'23_May_2008',  
granularity=>'auto');
```

```
exec dbms_stats.gather_table_stats('user',  
'emp', GRANULARITY => 'PARTITION');
```

```
exec dbms_stats.gather_table_stats('user',  
'emp', GRANULARITY => 'all');
```

*BEGIN*

*DBMS\_STATS.delete\_table\_stats ('winner','emp');*

*DBMS\_STATS.lock\_table\_stats ('winner','emp');*

*END;*

*Inform Oracle to use to **one day old** statistics*

```
dbms_stats.restore_schema_stats  
(ownname=> 'winner', as_of_timestamp =>  
systimestamp - 1,  
force      => TRUE)
```



*Hints*

## *Hints*

*Syntax must be correct or the Hint will be ignored, and no error message is issued.*

*There is a 255 character limit to Hints.*

*When using an alias for a table in the statement, the alias needs to be in the Hint.*

# *Hints Look Like?*

*DELETE /\*+ hint [text] \*/*

*INSERT /\*+ hint [text] \*/*

*SELECT /\*+ hint [text] \*/*

*UPDATE /\*+ hint [text] \*/*

*SELECT /\*+ FULL(table\_name) \*/ \* from emp ;*

*SELECT /\*+ INDEX(table\_name index\_name1  
index\_name2...) \*/*

*ORDERED - Force the driving table*

*SELECT /\*+ ORDERED \*/ column1,  
column2..FROM table1, table2*

*ALL\_ROWS - best throughput.*

*Select /\*+ ALL\_ROWS \*/ .....*

*The ALL\_ROWS hint usually suppresses an index*

*FIRST\_ROWS - best response time.*

*Select /\*+ FIRST\_ROWS \*/ .....*

*The FIRST\_ROWS hint usually forces an index*

***/\*+ROWID\*/** Use ROWID access method*

***/\*+USE\_NOCACHE\*/***

*Don't put the data in the buffers*

***/\*+USE\_CACHE\*/***

*Don't put the data in the buffers*

```
select /*+ parallel(e, 4) parallel(b, 4) */  
e.ename,hiredate,b.comm  
from   emp e, bonus b  
where  e.ename =b.ename
```

# *Hints for Joins*

*Nested loop*

*Use\_nl*

*Sort merge*

*Use\_merge*

*Hash*

*Use\_hash*



## *When Oracle ignores hint*

*SELECT /\*+ APPEND \*/ ...*

*SELECT /\*+ FIRST\_ROWS \*/ zip\_code  
from emp group by zip\_code;*

*SELECT /\*+ USE\_HASH(X) \*/ \* FROM  
inv\_master WHERE inv\_type = 'G';*

```
SELECT /*+ INDEX_DESC(TM TEAM_PK) */  
player_name,coach, date_joined from roster  
JOIN team t USING (tm_id)  
where t.effdt = (select max(effdt) From team  
TM where tm.tm_id = t.tm_id)
```

```
SELECT /*+ INDEX(books xisbn) */ * FROM  
books b WHERE b.isbn = :ISBN;
```

**INDEX\_JOIN**

*informs to use an index join as an access path*

**INDEX\_FFS** *Does a fast full index scan*

**INDEX\_SS** *Does an index skip scan*

**NO\_INDEX**

*Does not allow index scan and does full table scan*

*USE\_NL*

*Joins the specified table using a nested loop join*

*NO\_USE\_NL*

*Does not use nested loops to perform the join*

**USE\_HASH**

*Joins the specified table using a hash join*

**NO\_USE\_HASH**

*Does not use HASH JOIN to perform the join*

**USE\_MERGE**

*Joins the specified table using a merge join*

**NO\_USE\_MERGE**

*Does not use MERGE JOIN to perform the join*

*Append*

*Direct path insert*

*NOAPPEND*

*Regular insert*

*Parallel processing*




## *The Parallel Query Option*

*Multiple Server Processes can work together*


*Improves data-intensive and CPU intensive  
operations.*

*Available only when a Full table scan or a Sort  
operation is being performed.*

# Benefits of Parallel Execution

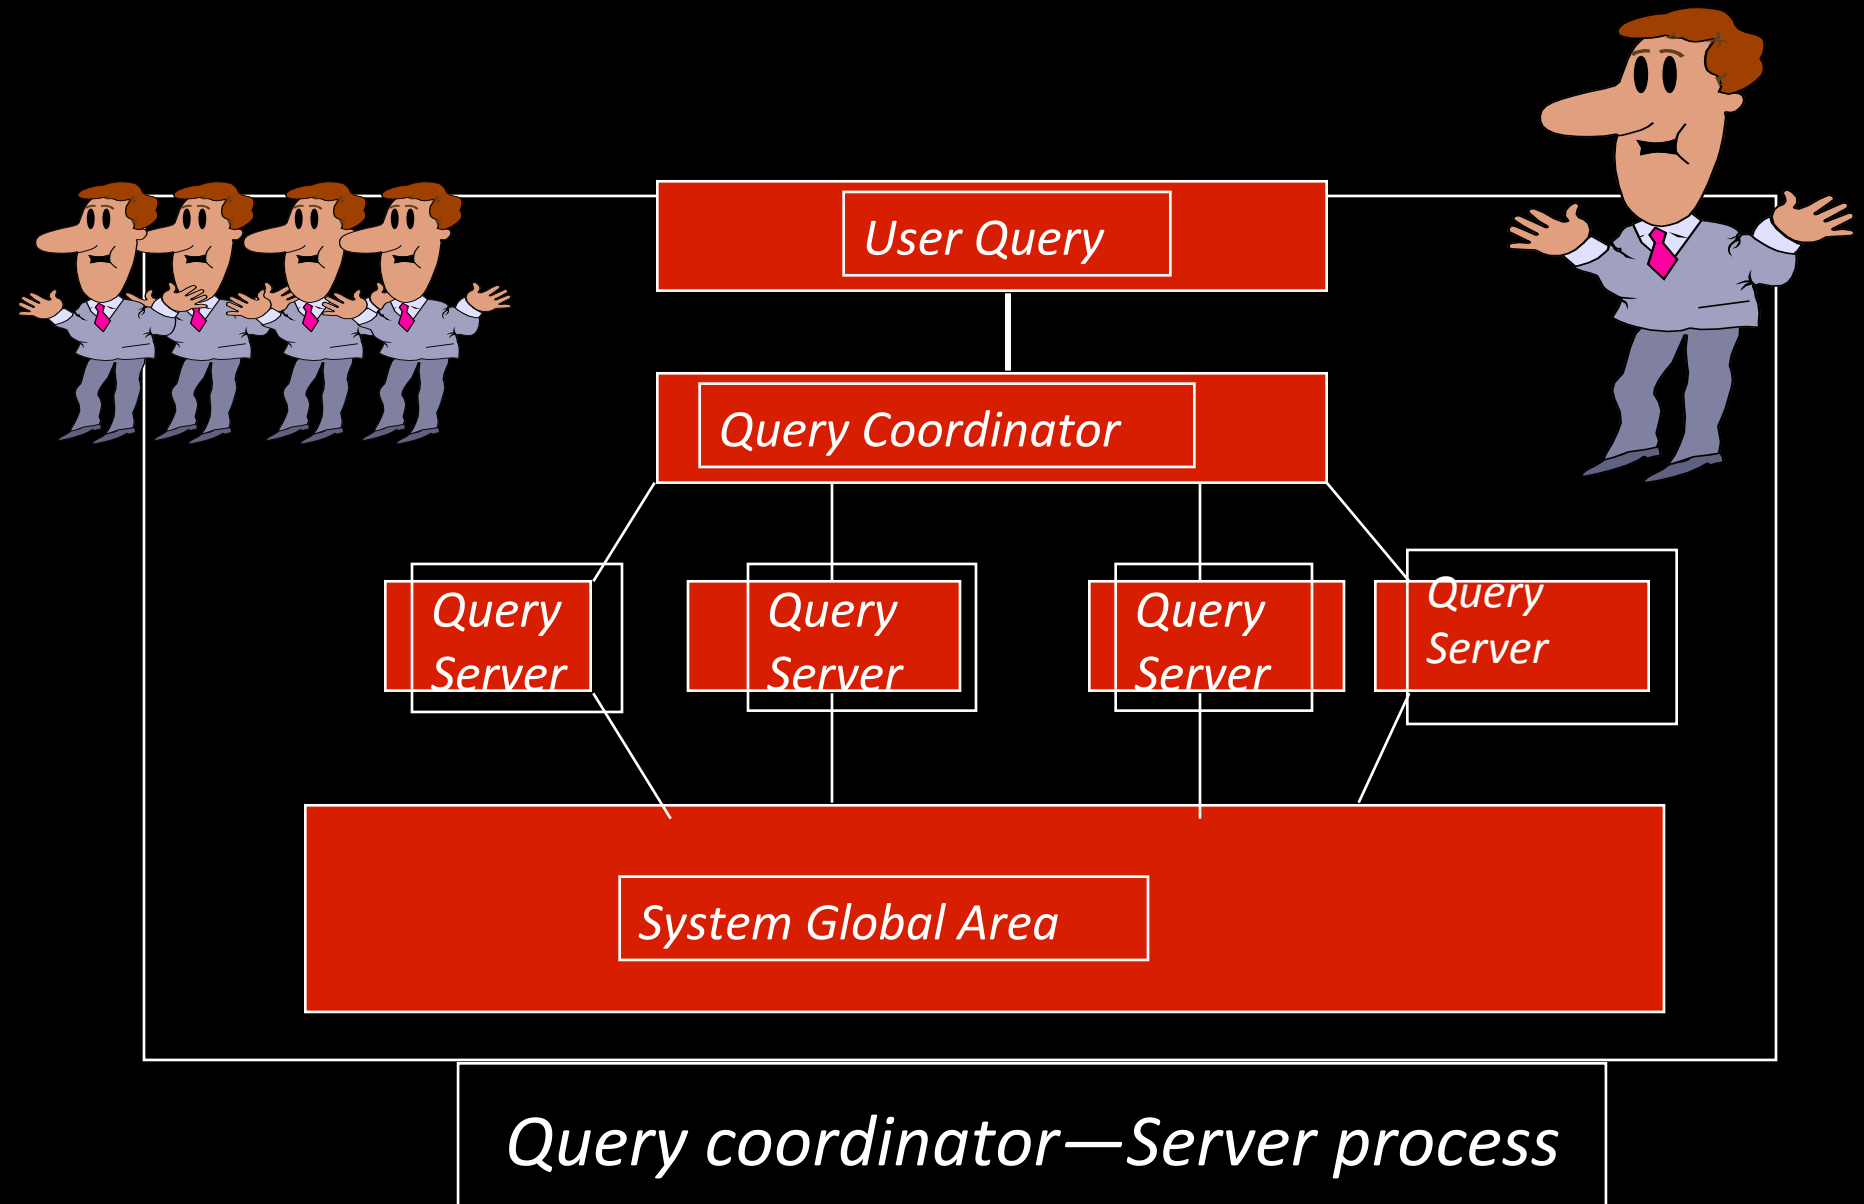


Server without Parallelism



Server with Parallelism

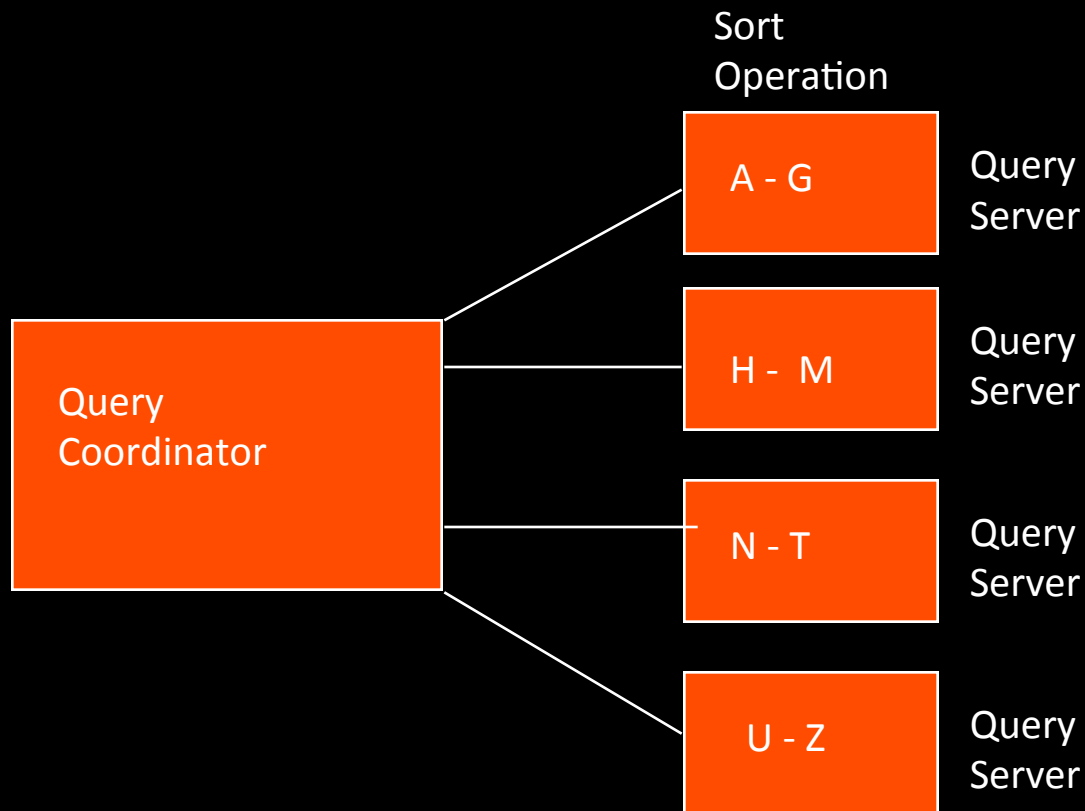
# *Parallel Query Process*



# Degree of Parallelism

Select degree from user\_tables where table\_name='EMP'

An example of a sort with the degree of parallelism set to 4



# Parallel Query Operations

*Access methods:*

*Index Range Scans*

*Various SQL operations:*

*Group by ,Order by, Not in, Exists, In,  
Distinct, Union,Union all, MinusTable  
Scans, Index Full Scans  
Partitioned ,Intersect,Cube, Rollup,  
Aggregates*

*Join methods:*

*Nested loop, Sort Merge, Hash,  
Star Transformation,  
partition-wise join*

## Parallel DDL Operations

*The parallel DDL statements for non-partitioned tables and indexes are:*

*CREATE INDEX*

*CREATE TABLE ...AS SELECT*

*ALTER INDEX ... REBUILD*

*The parallel DDL statements for  
partitioned tables and indexes are:*

*CREATE INDEX*

*CREATE TABLE ...AS SELECT*

*ALTER TABLE ... MOVE PARTITION*

*ALTER TABLE ... SPLIT PARTITION*

*ALTER TABLE ... COALESCE  
PARTITION*

*ALTER INDEX ... REBUILD  
PARTITION*

*ALTER INDEX ... SPLIT PARTITION*



## *Parallel Query Use*

### *With Hints*

```
Select /*+ Full(table) Parallel(table 3) */
```

### *Creating tables to use the Parallel Option*

```
Create Table emp (eno Number, tname  
Varchar2(10))
```

```
Parallel 2 ;
```

## Parallel DML - V8

Example (Parallel in subquery):

```
insert /*+ PARALLEL (t1,2) */ into t1  
(select /*+ PARALLEL (t2, 6) */ c1, c2, sum  
(c3) from t2 group by c1, c2);
```

*If you don't need parallel*

*Alter table emp noparallel ;*

*Select /\*+ noparallel \*/ \* from emp ;*

*JOINS*

*A join defines the relationship between two row sources.*

*A join is a method of combining data from two data sources.*

*It is controlled by join predicates, which define how the objects are related.*

*Nested loop*

*Sort\_merge*

*Hash join*

*Nested loop join*

*Cost of accessing outer table*

*+*

*(cardinality of outer table*

*\**

*cost of accessing inner table)*

*Sort merge join*

*(Cost of accessing outer table*

*+*

*outer sort cost)*

*+*

*(Cost of accessing Inner table*

*+*

*Inner sort cost)*



*Hash join*

*(Cost of accessing outer table)*

+

*(Cost of building hash table)*

+

*(Cost of accessing Inner table)*

# *What is a Driving Table?*

- Oracle refers to the driving table as the outer table in a join*
- This is the table that dictates how the data is retrieved from the database*

*Select T1.A, T1.B, T2.X, T2.Y from T1, T2*

*Where T1.A = 10 and*

*T1.B = T2.X;*

*T1-----Driving or Outer table*

*T2-----Driven or Inner table*


# *Which Table Should be the Driving Table?*

- The table that returns the fewest number of rows the fastest*
- If there are three or more tables, the driving table should be the intersection table*

# *Nested loop join*

*Oracle compares each row of an inner set  
with Each row of the outer set  
and returns those rows That satisfy the condition*


# Nested Loop Algorithms ...



# Hash Join ...Algorithms ...

- *The two tables are read and split into partitions.*
- *A hash table is built for each partition that fits into memory.*
- *Partitions that don't fit into memory are placed onto disk.*
- *The join is performed by taking a partition from the second table and probing the hash table.*

# Hash Join Diagram



The first table is hashed in memory, the second table is used to probe the hash (build) table for matches. In simple cases the cost is easy to calculate.



# *Hash Join ... Tips*

*HASH\_JOIN\_ENABLED =TRUE.*

*HASH\_AREA\_SIZE*

*Increase SORT\_AREA\_SIZE*

*Requires More Cpu*


# Sort / Merge Algorithms ...

*1. select count(t1.v1),count(t2.v2) from big1 t1,  
big2 t2 Where t2.n2 = t1.n1;*

*Oracle reads, filters, and usually sorts, the qualifying rows in each of the tables independently.*

*The two intermediate result tables are sorted in the same sequence based on the join predicate column(s).*

# Sort-Merge Join Diagram



# *Sort merge Join ... Tips*

*Larger tables that don't generate small result sets.*

*When indexes don't exist on the join predicates.*

*Requires more Sort area size and Temp segments*

*Require more Memory and Disk I/O*

*Basics*

*TUNING*

**Time**



**Money**



*“ Speed kills competitor ”*



**TARGET**

**User  
Satisfaction**





*Problem    Difference in expectation and reality*

*Satisfaction    --    Small difference*

*CHECK*

*User Experiences and Business Expectations*

# WHAT IS PERFORMANCE

*It is a measure of the responsiveness of a system*

# WHAT IS *PERFORMANCE* PROBLEM

*It is a lack of response within an expected time frame.*

*Performance improvement means*

*doing the same work in less time.*

*WHICH HAVE TO BE FASTER*

*Which have to be faster*

*1. User Response time problems*

*2. Inefficiencies that aren't yet noticeable  
as user response time problems.*

# Which requires Tuning

- Is business critical
- Executes frequently
- Runs a long time
- Consumes more resources



# WHEN TO TUNE

- *Application design and programming*
- *Database configuration*
- *Adding a new application*
- *Ongoing Tuning*

*HOW MUCH TO TUNE*

# **GOAL MUST BE CLEARLY DEFINED**

**QUERY X RUNS 30 MINUTES AND  
NEED TO RUN IN LESS THAN 2  
MINUTES**

**ONCE YOU**

**TARGET**

*HOW TO TUNE*

*ELIMINATION OF WORK STEPS :*

*REDUCING SERVICE TIME*

*ELIMINATION OF WAITERS*

*REDUCTION OF QUEUING DELAY*

*Causes of performance problems*

*1 : OVERLOAD*

*2 : YOU ARE DOING IT WRONG*



*What is my Target*

# *Method*

*Where is the time spent?*

*Identification*

*How is the time spent?*

*How to reduce the time spent?*

*Tuning*

# *Causes of CPU Usage*

- *Logical I/O*
- *Parsing*
- *Spinning (active waiting)*

## *Causes of Waiting*

- *Physical I/O*
- *Lock contention*
- *Latch contention*

*Have a Target (Reducing consistent gets) \*\*\**

*Design and develop with performance in mind*

*Use Bulk Methods*

*Reduce parsing \*\*\**

\*\*\* 100 % important

*Don't look for short cuts, keep it simple*

*Keep clear records of what has been changed*

*Have a sound reason for changing something*

*Change one thing at a time (Newton's law)*

*Keep proper plans for reverting back*

- *Use Proper table types*
- *Use Proper index types*
- *Use Parallel Processing*
- *Use proper Partition methods*
- *Use Analytic functions*

*Build your own test cases*

*Don't do redundant work*

*Know what database is capable of doing*

*Don't tune a query, tune a question*



*Definitions which may help you for the interview*

*What is explain plan*

*Text or graphical representation of  
SQL execution flow*

# ***Data Skew***

*A non-uniform distribution of data, generally on a per column, per value basis*

**Row Migration** - When the amount of data for a record exceeds the space within the block, the row is migrated to another block

**Solution:**

Increase the *PCTFREE* for the segment

**Row Chaining** - When the amount of data for a record exceeds the block size, the row is chained with another block

**Solution:**

*Increase the block size*

# Selectivity ( $P$ )

$$P = r / R$$

$R$  = The total number of rows in the table ( $R$ ):

$r$  = number of rows returned by the filter in  
where clause

# *Cardinality*

*Selectivity \* no of rows*

# *USER\_INDEXES*

*BLEVEL: Height of index between root block and leaf block*

*LEAF\_BLOCKS: Number of leaf blocks in index*

*DISTINCT\_KEYS: Number of distinct index values*

*AVG\_LEAF\_BLOCKS\_PER\_KEY: Average number of leaf blocks required to store an indexed value.*



## *USER\_INDEXES*

*AVG\_DATA\_BLOCKS\_PER\_KEY: Average number of table blocks that contain rows referenced by indexed key value*

*NUM\_ROWS: Number of leaf row entries*

*CLUSTERING\_FACTOR:*

*Indicates how well ordered the rows in the table are in relation to the index*

## *Fast full index scan –*

*This is an alternative to a full table scan when the index contains all the columns that are needed for the query.*

*At least one column in the index key has the NOT NULL constraint..*

*Uses multi block reads, unlike a full index scan.*

## *Full Index scan —*

*processes all of the leaf blocks of an index, but only enough of the branch blocks to find the first leaf block.*

*It uses single block reads.*

# *Parsing*

## *Syntactic check*

*Syntax, keywords*

## *Semantic check*

*Whether objects referenced exist, are accessible  
(by permissions) and are usable (status)*

*How to handle single user problem*

1. Understand the problem/experience from the user
2. Note down the current response time of the query
3. Understand user expectations ( from x seconds to y seconds )
4. Check with the user was there any increase /decrease in data volume
5. Check with the user was there any increase /decrease in user volume
6. Check with the user any changes were made in the query / Logic
7. Check whether statistics are gathered till date
8. Answer yourself for our 3 questions (use explain plan , v\$views)
9. Check reparsing , consistent gets
10. Write down solutions / benefits /efforts you will make for this issue
11. Map the right solution to the problem (Take more care here )
12. Do have a sound reason for changing some thing .
13. Cross check your solution will not bring issues in another area.
14. Have rollback plan if any thing goes wrong after your solution
15. Document all that you did for this issue .
  1. (For your reference)
  2. (You have a proof if anyone questions you )

*Thanks for your support*

*Thanks for your patience*

*Thanks for your Love*

*I respect and loved Your Hard work*

*I will never forget this crowd in life time*

*DONE*